

COL757 Programming Assignment

Yash Gupta – 2013CS10302

Q1: COMPUTE OR OF A MILLION BITS

- N = 1 mil => inconclusive results due to very small problem size.
- N = 10 mil

Threads/Algorithm	P = 6	P = 12	P = 24
SERIAL	0.007747160	0.007554330	0.007580700
CONCURRENT WRITE	0.000971427	0.000561381	0.000360807
BINARY RREE	0.000969006	0.000497586	0.000307906

- SERIAL runs in ~ constant time as expected
- Both parallel implementations give significant speedup compared to SERIAL, but only when problem size is large enough .
(N = 1 mil gives approximately same results for all algo/thread combinations)
- CONCURRENT WRITE is as fast as BINARY TREE when number of threads is low, but becomes increasingly less efficient with more threads.
- Conclusion: when operating on low number of threads, concurrent write can be useful due to its low implementation cost (in terms of software)

Q2: IMPLEMENT PARALLEL MERGESORT AND PARTITION SORT

- N = 16,000

Threads/Algorithm	P = 6	P = 12	P = 24
SERIAL	0.000990361	0.000993135	0.000997214
ODD-EVEN MERGE	0.00189207	0.00571344	0.0367163
PARTITION SORT	0.00814808	0.00948329	0.0214106

- N = 1,600,000

Threads/Algorithm	P = 6	P = 12	P = 24
SERIAL	0.142126	0.142126	0.142309
ODD-EVEN MERGE	0.0418275	0.0271082	0.0468103
PARTITION SORT	0.836293	1.0389	1.21488

- SERIAL performs better than both parallel implementations at smaller problem size
- ODD-EVEN MERGE outperforms SERIAL at large problem sizes.
- ODD-EVEN MERGE scales very moderately at large problem sizes, with degradation at P=24.
 - This might be due to the fact that there are only 12 physical cores on the machine, and 24 logical cores due to hyper-threading (super-scalar pipeline design).
- As thread count increases at smaller problem sizes, partition sort degrades moderately.
 - More rapid degradation is seen with large problem size (possibly due to large number (\sqrt{N}) of sub-problems created at each step)
- Conclusion:
 - For small problem sizes, use SERIAL
 - For larger problem sizes, use ODD-EVEN MERGE, as PARTITION SORT has very high overheads.